

---

# A Survey of No-Restoring Algorithm and Asynchronous Divider in RISC-V Processors for Low-Power Computing

---

[www.surveyx.cn](http://www.surveyx.cn)

## Abstract

This survey paper explores the integration of no-restoring algorithms and asynchronous dividers within RISC-V processors, highlighting their role in advancing low-power computing. The no-restoring algorithm enhances computational efficiency by streamlining division operations, reducing cycle counts, and minimizing power consumption, which is crucial in energy-constrained environments. Asynchronous dividers complement this by enabling non-blocking operations that improve processing speed and efficiency, particularly in heterogeneous and multicore systems. RISC-V's open-source and modular architecture facilitates extensive customization, allowing for the development of specialized instructions and extensions that optimize performance across various applications, including AI and machine learning. The incorporation of SIMD adaptations further boosts computational efficiency through parallel processing, essential for optimizing modular arithmetic. Future research should focus on exploring additional RISC-V extensions, optimizing compiler strategies, and enhancing the micro-decoding unit's flexibility. The findings suggest that RISC-V's adaptability presents opportunities for hardware optimizations and confirms its viability for complex software applications. Overall, the strategic integration of these algorithms and dividers within RISC-V processors represents a significant advancement in achieving high-performance, energy-efficient computing, with numerous avenues for future exploration and development.

## 1 Introduction

### 1.1 Importance of Low-Power Computing

Low-power computing is essential in modern technology, driven by the demand for energy-efficient solutions across various applications, particularly in edge computing, where efficient arithmetic is critical for performance and energy savings [1]. The transition to resource-constrained hardware, especially in the Internet of Things (IoT), emphasizes energy efficiency for low-power smart devices [2]. As the market for small, battery-powered IoT endpoint devices expands, the need for ultra-low power operation and flexible computing capabilities becomes increasingly important [3].

Challenges in enhancing processor performance due to the decline in lithographic scaling further underscore the significance of low-power computing for sustaining technological advancement [4]. The demand for efficient data processing in edge computing applications highlights the impact of low-power computing on energy efficiency [5]. Additionally, enabling On-Device Learning (ODL) for ultra-low-power Micro-Controller Units (MCUs) necessitates energy-efficient solutions, further emphasizing the broader relevance of low-power computing [6].

The increasing complexity of IoT devices necessitates processors that can balance performance, efficiency, and cost [7]. As Moore's law nears its limits, the industry must explore alternatives

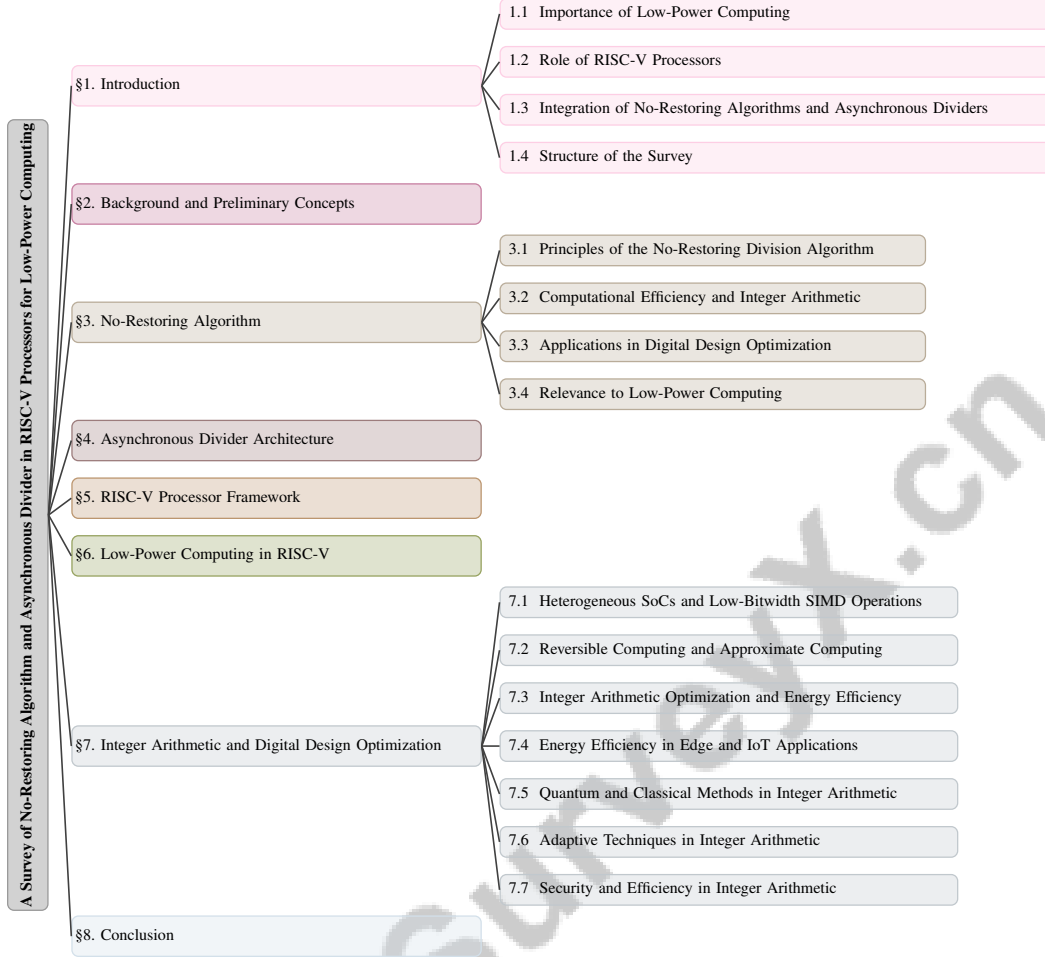


Figure 1: chapter structure

to miniaturization to meet the rising computational demands in embedded systems [8]. The need for efficient computation in edge devices, particularly for large language model inference, further highlights the critical role of low-power computing in contemporary technology [9]. Collectively, these factors illustrate the vital importance of low-power computing in advancing energy efficiency and sustainability across various technological domains.

## 1.2 Role of RISC-V Processors

RISC-V processors significantly advance low-power computing through their open-source and modular architecture, which fosters customization and innovation across applications [10]. The open Instruction Set Architecture (ISA) of RISC-V allows developers to adapt the ISA to meet specific application needs, enhancing computational efficiency and reducing power consumption [9]. This adaptability is crucial for deploying lightweight neural networks and other machine learning applications on RISC-V-based processors, as demonstrated by the FANN-on-MCU toolkit [2].

The modularity of the RISC-V ISA enables the incorporation of instruction extensions and microarchitectural optimizations, facilitating near-threshold operation essential for energy efficiency in edge computing environments [10]. This flexibility allows RISC-V processors to effectively support diverse application domains, including IoT and System-on-Chip (SoC) applications, while providing robust security features and an open-source framework [11]. Moreover, RISC-V's involvement in secure remote attestation frameworks, such as SRA-Care, highlights its capability for enabling lightweight secure communication, critical for low-power computing [12].

---

In Multi-Processor Systems-on-Chip (MPSoCs), open-source RISC-V cores are increasingly used for efficient local data processing, underscoring their importance in IoT applications [13]. The integration of RISC-V processors into systems like SentryCore demonstrates their reliability for real-time subsystems, which can be seamlessly integrated into mixed-criticality systems [14]. Furthermore, the development of cost-effective soft processors on FPGAs using RISC-V illustrates its potential for scalable and flexible computing solutions [15].

The open-source and modular nature of RISC-V processors, combined with their ability to integrate specialized functionalities, positions them as a cornerstone in the evolution of low-power computing, addressing the increasing demands for energy-efficient and high-performance applications [8].

### 1.3 Integration of No-Restoring Algorithms and Asynchronous Dividers

Integrating no-restoring algorithms and asynchronous dividers within RISC-V processors enhances computational efficiency, especially in energy-constrained environments. No-restoring division algorithms streamline division operations by omitting the iterative correction steps found in traditional restoring methods, thus reducing computational overhead. This approach aligns with the energy-efficient design philosophy of RISC-V processors, which prioritize minimizing power consumption while maintaining performance [4]. By decreasing the cycles required for division, these algorithms contribute to faster computation times and lower energy usage, critical for low-power applications [1].

Asynchronous dividers complement this by enabling non-blocking operations that enhance overall processing speed and efficiency of RISC-V architectures. The asynchronous design allows computation and communication to overlap, decreasing idle times and increasing throughput. This advantage is particularly significant in heterogeneous computing environments where mixed-precision integer arithmetic is utilized, as demonstrated in frameworks like DARKSIDE, which leverage these capabilities for enhanced computational performance [16]. The integration of custom vector dot product instructions into the RISC-V architecture further exemplifies how tailored solutions can boost computational efficiency in edge AI applications [9].

Moreover, the synergy between no-restoring algorithms and asynchronous dividers supports deploying specialized accelerators within RISC-V processors, facilitating advanced applications such as neuromorphic computing [5]. The NoX processor, a proposed 32-bit RISC-V core, illustrates the potential of integrating application-specific hardware accelerators to enhance data processing efficiency in MPSoCs [13]. Additionally, incorporating secure communication frameworks like SRA-Care underscores the importance of combining computational efficiency with robust security measures in modern processor design [12].

The integration of no-restoring algorithms and asynchronous dividers within RISC-V processors represents a strategic approach to achieving high-performance, low-power computing. This integration addresses the computational demands of contemporary applications while aligning with the broader objectives of energy efficiency and sustainability in digital design [3].

### 1.4 Structure of the Survey

This survey provides a comprehensive analysis of integrating no-restoring algorithms and asynchronous dividers within RISC-V processors, emphasizing their role in enhancing low-power computing. It begins with an introduction that outlines the critical importance of low-power computing in modern technology and the essential role of RISC-V processors in advancing this field. The discussion highlights how RISC-V's open-source architecture fosters innovation in high-performance and low-power applications, addressing the challenges and opportunities presented by integrating RISC-V processors into heterogeneous, multicore system-on-chip (SoC) designs, particularly concerning artificial intelligence and machine learning workloads [17, 18, 19]. The survey then delves into the integration of no-restoring algorithms and asynchronous dividers, underscoring their contributions to computational efficiency.

Following the introduction, Section 2 provides background on preliminary concepts, including the modularity of RISC-V architecture and principles of low-power computing. It explores integer arithmetic and digital design optimization, laying the foundation for understanding subsequent sections.

---

Section 3 focuses on the no-restoring algorithm, discussing its principles, computational efficiency, and applications in digital design optimization. The significance of this algorithm in low-power computing is explored, particularly in light of recent advancements in computer arithmetic that prioritize energy efficiency and performance optimization for edge computing applications, such as those in the IoT. These advancements include adopting novel data types and techniques, such as bfloat16 and fixed-point arithmetic, crucial for meeting the power and area constraints of embedded systems while delivering necessary computational capabilities for machine learning and artificial intelligence tasks [19, 1].

Section 4 thoroughly examines the architecture of the asynchronous divider, focusing on its design principles and the resulting effects on performance metrics and power efficiency, particularly in multicore SoC implementations utilizing RISC-V processors. This exploration highlights how integrating such dividers can enhance computational capabilities while addressing stringent power constraints often encountered in edge computing and IoT applications [20, 1, 21, 10, 18]. The adaptability of asynchronous dividers to SIMD architectures and their applications in heterogeneous and multicore systems are also discussed.

Section 5 delves into the RISC-V processor framework, highlighting its open-source nature and modular design, which fosters innovation across diverse applications—from high-performance to low-power cores—and facilitates integrating heterogeneous multicore SoC architectures. This section discusses how the RISC-V ecosystem, supported by a growing number of open-source implementations, addresses critical design challenges while promoting energy efficiency, security, and adaptability in modern computing environments [17, 18]. It covers integrating micro-decoding units, instruction set extensions, and the importance of formal verification and validation.

Section 6 explores various strategies for achieving low-power computing in RISC-V processors, emphasizing integrating specialized accelerators and dynamic instruction execution. It highlights the benefits of incorporating a dynamic micro-decoder unit to enable custom instruction sequences, optimizing tasks such as binary compression and behavior obfuscation. Additionally, this section discusses developing a specialized instruction set for edge AI applications, designed to enhance execution efficiency of large language model (LLM) computations while significantly reducing energy consumption. This approach meets the demands of edge devices and illustrates the potential of RISC-V architectures in facilitating heterogeneous, multicore SoC designs that can efficiently support multithreaded applications [17, 9, 18]. The trade-offs between power efficiency and computational performance are critically analyzed.

In Section 7, the role of integer arithmetic in digital design optimization is analyzed, discussing techniques for optimizing digital design to enhance energy efficiency and processing speed. This section also addresses security and efficiency in integer arithmetic operations.

The conclusion summarizes the survey’s key findings, emphasizing the benefits of integrating no-restoring algorithms and asynchronous dividers in RISC-V processors for low-power computing. It highlights potential areas for future research and development, aligning with the objectives of energy efficiency and sustainability in digital design [14]. The following sections are organized as shown in Figure 1.

## **2 Background and Preliminary Concepts**

### **2.1 RISC-V Architecture and Modularity**

The open-source and modular nature of the RISC-V architecture has catalyzed its widespread adoption across various domains, from embedded systems to high-performance computing. This flexibility allows developers to tailor and extend the architecture to suit specific requirements, fostering innovation [13]. RISC-V supports a range of instruction set extensions, such as RV32I, which is foundational for operating systems and well-suited for embedded applications [15].

The architecture’s adaptability is further demonstrated through specialized cores and peripherals, including the RISC-V RV32IM core and a PLIC-based interrupt controller, both designed for extensibility and configurability, which are crucial for adaptable virtual prototypes [7]. The integration of SentryCore into Systems-on-Chip via AXI4 interfaces exemplifies RISC-V’s modularity, facilitating seamless system integration [14].

---

RISC-V also permits the creation of custom instructions tailored for specific applications, notably in vector dot product calculations that are vital for accelerating large language models [9]. The STRELA architecture, a coarse-grained reconfigurable architecture, illustrates RISC-V's capability to efficiently manage both control-driven and data-driven applications via its elastic logic and independent memory nodes [8].

Benchmarking open-source RISC-V implementations has standardized comparisons of performance, resource utilization, and power consumption, guiding developers in selecting optimal implementations. The evaluation of NoX against other resource-constrained RISC-V cores highlights its compact architecture and integration ease, reinforcing the advantages of RISC-V's modularity [13].

RISC-V's modularity and open-source philosophy position it as a leader in processor design, enabling the development of customizable, energy-efficient processors that meet the evolving demands of modern applications. This adaptability enhances performance across tasks, including artificial intelligence and machine learning, while addressing critical challenges in security and power efficiency. The expanding ecosystem of RISC-V cores and features like dynamic micro-decoding pave the way for innovative architectures that support multicore and heterogeneous systems, driving advancements in research and commercial applications [17, 18, 19].

## 2.2 Low-Power Computing and Energy Efficiency

Enhancing low-power computing and energy efficiency is vital in modern computing architectures, particularly for RISC-V processors. The increasing demand for energy-efficient solutions is driven by machine learning applications that require reliable, power-efficient architectures. Utilizing integer arithmetic is one strategy for achieving low power consumption, avoiding overflow while maintaining accuracy in sparse linear algebra computations [4]. This aligns with optimizing computational processes to minimize power usage.

RISC-V architectures are well-suited for addressing low-power computing challenges due to their modular and open-source nature, which facilitates extensive customization and optimization of hardware resources [10]. The integration of modular integer arithmetic and reduced precision techniques significantly enhances energy efficiency, particularly in optimizing training primitives for deep neural networks (DNNs) on resource-constrained microcontroller units (MCUs) [6]. This is crucial for meeting the stringent demands of TinyML applications, which require high performance with low power consumption.

Efficient on-chip DNN inference and training are essential for low-power smart devices, necessitating careful management of power and area budgets to achieve desired performance levels. Benchmarking various RISC-V implementations highlights challenges in direct performance comparisons due to configuration and technology variations, emphasizing the need for standardized evaluation approaches for energy efficiency across architectures [7]. Neuromorphic architectures have been benchmarked for speed, energy consumption, and footprint, facilitating comparisons and showcasing their potential for energy-efficient computing [5].

Near-threshold RISC-V cores address inefficiencies in managing memory hierarchies and computational demands, improving energy consumption and performance in data-intensive workloads [10]. These advancements are essential for meeting the increasing demands for energy-efficient computing across various technological domains, aligning with sustainability and performance optimization objectives in digital design.

## 2.3 Integer Arithmetic in RISC-V Processors

Integer arithmetic is crucial for the functionality and performance optimization of RISC-V processors, significantly influencing digital design strategies. The open-source nature of RISC-V facilitates the integration of specialized integer arithmetic operations that enhance computational efficiency and reduce power consumption in embedded systems [20]. Benchmark programs like Dhrystone, Coremark, and Embench are commonly used to assess processor performance, underscoring the importance of efficient integer arithmetic in both academic and industrial contexts.

The efficiency of integer arithmetic is highlighted by challenges in classical methods, particularly in multiply-add operations, which often face depth complexity constraints when executed in sequence,

---

leading to inefficiencies [22]. RISC-V processors mitigate these inefficiencies through optimized integer arithmetic circuits that minimize resource usage and enhance processing speed.

In digital design, integer arithmetic specifications must be ground-complete and intuitive to facilitate efficient computation and representation. This requirement is critical for developing effective digital systems capable of handling complex arithmetic operations with minimal overhead [23]. RISC-V supports this by enabling customization of arithmetic operations to meet specific application needs, optimizing both performance and energy efficiency.

Moreover, the satisfiability of formulas in Integer Arithmetic, particularly under Satisfiability Modulo Theories (SMT), is a significant consideration in digital design. Accurately determining the satisfiability of Boolean combinations of arithmetic formulas and propositional variables is essential for reliable execution of integer arithmetic operations in RISC-V processors [24]. This capability is vital for supporting advanced computational models and ensuring the reliability of digital systems.

The integration of integer arithmetic in RISC-V processors also addresses limitations in probabilistic programming languages concerning integer distributions. Efficient handling of integer arithmetic is crucial for inference on complex models involving intricate arithmetic operations [25]. By optimizing integer arithmetic, RISC-V processors can effectively support a broad spectrum of applications, from embedded systems to high-performance computing, thereby enhancing the overall digital design landscape.

## 2.4 Digital Design Optimization Techniques

Digital design optimization techniques are essential for enhancing energy efficiency and processing speed in computing systems, particularly within RISC-V processors. These techniques often integrate advanced computational models and architectural innovations that address traditional digital design limitations. One notable approach is the implementation of hybrid quantum circuits, such as the Quantum Multiply-Add Circuit (QMAC), which utilizes the Quantum Fourier Transform (QFT) to perform integer arithmetic operations more efficiently than classical circuits, significantly optimizing digital design by reducing computational overhead [22].

Reversible computing presents another promising avenue for optimization. By employing reversible computing principles through category theory, including concepts like dagger categories and inverse categories, designers can create systems that minimize energy dissipation during computation. This approach enhances energy efficiency and aligns with sustainable computing practices by reducing the thermal footprint of digital systems [26].

In integer arithmetic, a systematic approach to defining datatype defining rewrite systems (DDRSes) ensures ground-completeness and practical applicability of integer operations. This methodology enhances digital system efficiency by providing a robust framework for managing complex arithmetic operations with minimal resource utilization [23]. Additionally, representing integer distributions as distributions on binary encodings allows for more efficient manipulation and inference, overcoming inefficiencies associated with traditional one-hot categorical encoding [25].

The adaptability of digital design architectures, exemplified by the STRELA architecture, further showcases optimization potential in energy efficiency and processing speed. STRELA's capacity to execute both control-driven and data-driven applications through elastic logic and independent memory nodes underscores its superiority over traditional Coarse-Grained Reconfigurable Architectures (CGRAs), enhancing overall energy efficiency and performance [8].

Recent research emphasizes the critical need for integrating advanced computational models and architectural strategies, such as next-generation computer arithmetic and RISC-V-based hardware accelerators, to improve performance and energy efficiency in contemporary computing systems, particularly in resource-constrained environments like edge computing and artificial intelligence applications [19, 1, 27, 9, 18]. By leveraging advancements in quantum computing, reversible computing, and adaptable architectures, RISC-V processors can achieve substantial improvements in processing speed and energy consumption, meeting the growing demands of modern digital applications.

Category	Feature	Method
<b>Principles of the No-Restoring Division Algorithm</b>	Arithmetic Optimization	ACMQFT[28], ECLT[29]
<b>Computational Efficiency and Integer Arithmetic</b>	Integer Arithmetic Optimization	QMAC[22], IAOCNN[30], RVP[15], RKM[31], LS[24], BEID[25]
<b>Applications in Digital Design Optimization</b>	Algorithm Optimization	SQED-SS[32]
<b>Relevance to Low-Power Computing</b>	Energy Optimization	STRELA[8], FANN-MCU[2], RFPF-ODL[6], SC[14]
	Integrated Design	VDOT[9]

Table 1: This table provides a comprehensive overview of various methods and features related to the no-restoring division algorithm. It categorizes the methods into principles of the algorithm, computational efficiency, applications in digital design optimization, and relevance to low-power computing, highlighting specific techniques and references to key studies. The table serves as a resource for understanding the diverse applications and optimizations associated with the no-restoring division algorithm in advancing computational efficiency and energy performance.

### 3 No-Restoring Algorithm

Understanding the no-restoring algorithm’s operational efficiency in various computational contexts is essential, especially given recent advancements in computer arithmetic that cater to the needs of low-cost, high-performance computations in edge computing and machine learning. Table 1 presents a detailed categorization of methods and features associated with the no-restoring division algorithm, emphasizing its principles, computational efficiency, applications in digital design, and relevance to low-power computing. This analysis highlights the algorithm’s alignment with emerging data types and techniques like mixed-precision arithmetic, which are increasingly relevant [19, 9, 23, 1]. The no-restoring division algorithm offers a robust framework for understanding its advantages over traditional methods, impacting digital arithmetic and processor design.

#### 3.1 Principles of the No-Restoring Division Algorithm

The no-restoring division algorithm is a pivotal method in digital arithmetic, streamlining division operations by eliminating iterative correction steps typical of traditional restoring methods. This approach enhances computational efficiency by reducing the cycle count, crucial for energy-constrained environments, aligning with the broader goal of optimizing computational processes for minimal power consumption [1]. In processor design, the implementation of no-restoring algorithms requires generating efficient instruction streams at the Register Transfer Level (RTL) to verify processor functionality [29]. The algorithm’s principles are similar to those in constructing quantum circuits for multilevel qudits, focusing on computational efficiency through innovative design [28].

The algorithm works by iteratively subtracting the divisor from the dividend, maintaining a running quotient tally until the dividend is less than the divisor, determining the remainder. This method bypasses restoring steps, facilitating faster computation and reduced complexity, making it ideal for application-specific and hardware-oriented designs [1]. Integrating no-restoring algorithms into RISC-V processors exemplifies enhanced computational efficiency through specialized arithmetic operations, as seen in advanced processing units like the Nanhv-vdot processor, which improves performance with dedicated vector dot product units [9]. Optimizing instruction fetch units and ALUs in processors like RVCoreP further emphasizes the importance of efficient arithmetic operations for high-performance computing [15].

As illustrated in Figure 2, the hierarchical structure of the no-restoring division algorithm highlights its efficiency in digital arithmetic, its implementation in processor design, and its applications across various computational environments. This foundational algorithm enhances division efficiency, which is critical across applications from embedded systems to high-performance computing. By optimizing these operations, the algorithm supports improving energy efficiency and computational performance in modern processor architectures, relevant to emerging trends in computer arithmetic like high-precision anchored numbers and mixed-precision computing, increasingly adopted in edge computing for low-cost, accurate arithmetic in machine learning and IoT applications [27, 1].

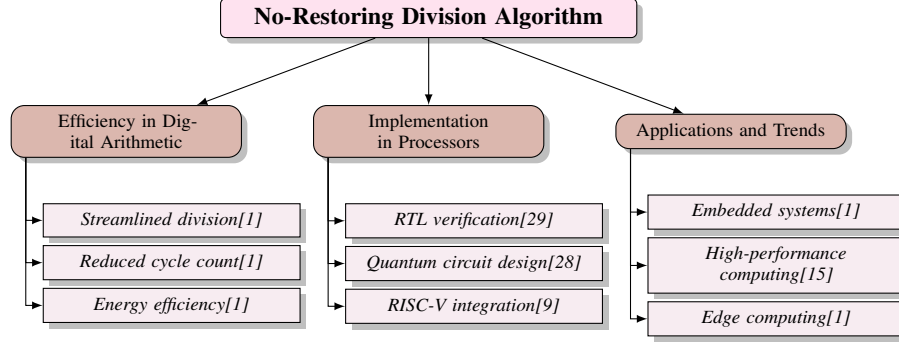


Figure 2: This figure illustrates the hierarchical structure of the no-restoring division algorithm, highlighting its efficiency in digital arithmetic, implementation in processor design, and applications in various computational environments.

### 3.2 Computational Efficiency and Integer Arithmetic

The no-restoring algorithm enhances computational efficiency in integer arithmetic by streamlining the division process and reducing the overhead of traditional methods. Its efficiency is evident in minimizing cycle counts for division, accelerating computation, and reducing power consumption, crucial for energy-constrained environments [30]. This aligns with optimizing digital arithmetic operations for high performance and energy efficiency.

In integer arithmetic, the no-restoring algorithm's impact is mirrored in innovative approaches enhancing computational efficiency, such as the Quantum Multiply-Add Circuit (QMAC), reducing circuit depth compared to classical methods, showcasing quantum computing's potential in arithmetic operations [22]. This mirrors the no-restoring algorithm's optimization of division operations, highlighting efficient arithmetic techniques' importance in modern computing architectures.

Integer arithmetic optimization is exemplified by the RVCoreP processor, streamlining critical processing paths to increase operating frequency and performance [15]. This is crucial for supporting advanced computational models and digital system reliability. Adaptive bounds on the BReLU activation function minimize quantization noise, ensuring effective integer network performance [30], underscoring precise arithmetic operations' significance in enhancing computational efficiency.

Advanced methods like LocalSMT utilize local search techniques for satisfiability modulo theories (SMT) problems in integer arithmetic, demonstrating superior performance over traditional methods [24]. This aligns with the no-restoring algorithm's focus on optimizing arithmetic operations for better computational efficiency. The asymptotic reduction of space complexity in Karatsuba-based multiplication circuits illustrates potential improvements in arithmetic circuit design [31].

Efficient binary encoding methods enhance probabilistic programming languages' performance and scalability in handling complex integer arithmetic [25]. This emphasizes optimizing integer arithmetic operations to achieve high computational efficiency, aligning with modern processor architectures' goals for energy-efficient and high-performance computing solutions.

### 3.3 Applications in Digital Design Optimization

The no-restoring division algorithm significantly enhances performance and energy efficiency in digital design optimization by streamlining arithmetic operations. This is crucial in resource-limited environments like embedded systems and IoT devices, where reduced division operation complexity and cycle count contribute to faster processing speeds and lower power consumption. This efficiency benefits digital systems relying on modular arithmetic, where SIMD (Single Instruction, Multiple Data) adaptations have shown substantial performance improvements [33].

Incorporating the no-restoring algorithm into digital design allows developing efficient arithmetic units handling complex operations with reduced latency. This is crucial for processors performing high-frequency arithmetic tasks, especially in edge computing, where low-power, high-performance solutions are increasingly demanded. Recent computer arithmetic advancements, like high-precision anchored numbers and posit arithmetic, offer alternatives to traditional IEEE 754-2008 compliant



arithmetic, optimizing performance in applications from embedded systems to high-performance computing. Specialized instruction sets in processors, like RISC-V-based ones, accelerate operations like vector dot products, significantly improving execution efficiency and reducing energy consumption in edge AI applications [9, 1]. By minimizing division operation overhead, the no-restoring algorithm enables digital systems to achieve higher throughput and better energy efficiency, crucial for modern applications.

The algorithm in digital design not only streamlines adaptive technique integration but also boosts system performance by leveraging advanced arithmetic methods like bfloat16 and mixed-precision computing, essential for optimizing power efficiency and computational accuracy in edge computing environments [9, 27, 1, 23]. Integrating specialized hardware accelerators leveraging the no-restoring algorithm optimizes arithmetic-intensive workload execution, improving overall processor efficiency. This aligns with digital design optimization objectives focusing on balancing computational performance and energy consumption.

stage (write-back stage) of the pipeline. At this time, all SIF instructions must have committed, as the pipeline is occupied by Symbolic QED instructions.

Specifies of the  $T_c$  recorder for a 5-stage, single-fetch, in-order pipeline is given in Fig. 5. Inputs are ready signals for all stages that precede the commit stage (e.g.,  $fetch\_ready$  is high when the fetch stage is ready to receive an instruction). The output  $SIF\_complete$  is true when the first Symbolic QED instruction goes through all pipeline stages and reaches the commit stage. The output  $mode$  keeps track of progress made so far by the Symbolic QED instruction (two later make use of this output in the Symbolic QED operand recorder). This  $T_c$  recorder for a 5-stage pipeline can be easily modified to support in-order pipelines with a different number of stages.

For an OoO core, the  $T_c$  recorder is even simpler, and utilizes the reorder buffer (ROB). The idea is to mark the entry allocated in the ROB for the first Symbolic QED instruction. After this,  $SIF\_complete$  is assigned true when the ROB head pointer reaches the marked instruction.


For cores with no ROB, but OoO commit (e.g., [Aquarius]), an additional constraint is required (see Sec. 2.3).

```


INPUT: fetch_ready, dispatch_ready, exec_ready, mem_ready
OUTPUT: SIF_complete, mode
// initialization
mode = S; SIF_complete = false;
// end initialization
if (mode == S) && (fetch_ready) then
mode = S; SIF_complete = false; // instruction completes fetch stage
end if
if (mode == S) && (dispatch_ready) then
mode = S; SIF_complete = false; // instruction completes decode stage
end if
if (mode == S) && (exec_ready) then
mode = S; SIF_complete = false; // instruction completes execute stage
end if
if (mode == S) && (mem_ready) then
mode = S; SIF_complete = true; // instruction completes memory stage
end if

```

(a) A C++ code snippet for a pipeline system[32]



(b) A complex diagram illustrating various mathematical operations and their interactions[1]



(c) Google Scholar hits and FPGA boards[21]

Figure 3: Examples of Applications in Digital Design Optimization

As shown in Figure 3, the No-Restoring Algorithm is pivotal in digital design optimization, offering significant insights into computational efficiency and hardware design. It is instrumental in applications demanding high-speed arithmetic operations, as depicted in the figures. The first image illustrates a C++ code snippet for a pipeline system, showcasing a five-stage, in-order processor pipeline optimizing data flow and processing speed. The second image presents a complex diagram of mathematical operations, emphasizing the algorithm's capability to handle intricate computational interactions, crucial for optimizing digital circuits. Lastly, the radar chart comparing Google Scholar hits and FPGA boards underscores the algorithm's impact and reach within the research community, revealing its adoption across various projects like Rocket, BOOM, and CVA6/Ariane. Together, these examples underscore the No-Restoring Algorithm's role in enhancing digital design through efficient computation and optimization strategies [32, 1, 21].

### 3.4 Relevance to Low-Power Computing

The no-restoring algorithm is crucial for advancing low-power computing goals, particularly within RISC-V processor architectures. By optimizing division operations and minimizing cycle counts, it significantly reduces computational overhead, lowering power consumption and enhancing performance in energy-constrained environments. Its integration into architectures like the RVCoreP-32IM facilitates efficient M-extension instruction execution, essential for high computational efficiency applications [20].

The algorithm's relevance is underscored by its application in edge AI environments, enhancing performance while aligning with low-power computing objectives [9]. Its support for advanced machine learning applications on resource-constrained devices without compromising performance is exemplified by the FANN-on-MCU toolkit, enabling efficient neural network deployment on ultra-low-power microcontrollers [2].

---

Additionally, the algorithm’s role in enhancing energy efficiency is demonstrated in architectures exploiting near-threshold voltage operation and parallelism, like certain RISC-V cores. These maintain computational performance while significantly improving energy efficiency, crucial for sustaining low-power computing demands [3]. The STRELA architecture further exemplifies this by achieving peak performance of 1.22 GOPs and energy efficiency of 115.96 MOPs/mW, demonstrating significant performance improvements and energy efficiency for embedded applications [8].

Moreover, integrating the algorithm into real-time on-device learning (ODL) on low-power devices supports efficient arithmetic operations without sacrificing accuracy, aligning with broader low-power computing objectives [6]. In safety-critical applications, architectures like SentryCore, featuring a triple-core lockstep design with ECC-protected memory, enhance reliability while supporting low-power computing goals [14]. The algorithm’s application within such architectures ensures efficient and reliable computational processes, essential for maintaining system integrity in energy-sensitive environments.

The no-restoring algorithm plays a crucial role in advancing low-power computing by optimizing arithmetic operations, reducing energy consumption, and enhancing performance across applications, particularly in resource-constrained environments like edge computing and IoT devices. It supports innovative arithmetic techniques like approximate computing and mixed-precision arithmetic, essential for high-performance applications while minimizing power usage [9, 6, 1, 13]. Its integration into modern processor architectures aligns with the pursuit of energy-efficient and high-performance computing solutions, meeting contemporary digital applications’ demands.

In recent years, the exploration of advanced computational architectures has garnered significant attention, particularly in the context of enhancing performance and power efficiency. A pivotal development in this domain is the Asynchronous Divider Architecture, which presents a novel approach to division operations in computing systems. This architecture not only addresses traditional performance bottlenecks but also optimizes energy consumption, making it particularly relevant for heterogeneous and multicore environments.

To illustrate these concepts, Figure 4 provides a comprehensive depiction of the hierarchical structure of the Asynchronous Divider Architecture. This figure details its design principles and highlights key features and optimization techniques that contribute to its effectiveness. Moreover, it outlines the performance benefits and efficiency enhancements offered by asynchronous dividers, thereby showcasing their adaptability and significant contributions to modern computing environments. By integrating these insights, we can better understand the transformative potential of asynchronous designs in advancing computational efficiency and energy performance.

## 4 Asynchronous Divider Architecture

### 4.1 Design Principles and Components

The asynchronous divider architecture is engineered to enhance computational efficiency and energy performance, particularly in power-constrained environments. By eliminating global clock distribution, this architecture leverages asynchronous design benefits, including reduced power consumption and improved speed through non-blocking operations that facilitate overlapping computation and communication, thereby minimizing idle times and maximizing throughput [16]. A key feature is the integration of mixed-precision integer arithmetic, crucial for efficient DNN processing, employing dedicated accelerators to optimize arithmetic operations and enhance the performance of RISC-V cores within heterogeneous computing clusters [16]. SIMD extensions further enable parallel data processing, maximizing energy efficiency in IoT applications [3].

As illustrated in Figure 5, the hierarchical structure of design principles and components in asynchronous divider architecture encompasses various memory optimization techniques and virtual prototypes. This figure highlights key elements such as efficiency enhancement, mixed-precision arithmetic, SIMD extensions, and the role of SystemC and TLM-2.0 in early software development and system-level validation. Advanced memory access patterns and a smart L0 buffer optimize data handling and reduce latency, maintaining high performance while minimizing energy consumption in resource-constrained settings [3]. Transforming training kernels into Matrix Multiplication operators and using FP16 SIMD instructions align with asynchronous architecture design principles, enhancing performance [6]. The RVCoreP method exemplifies optimizations in instruction fetch,

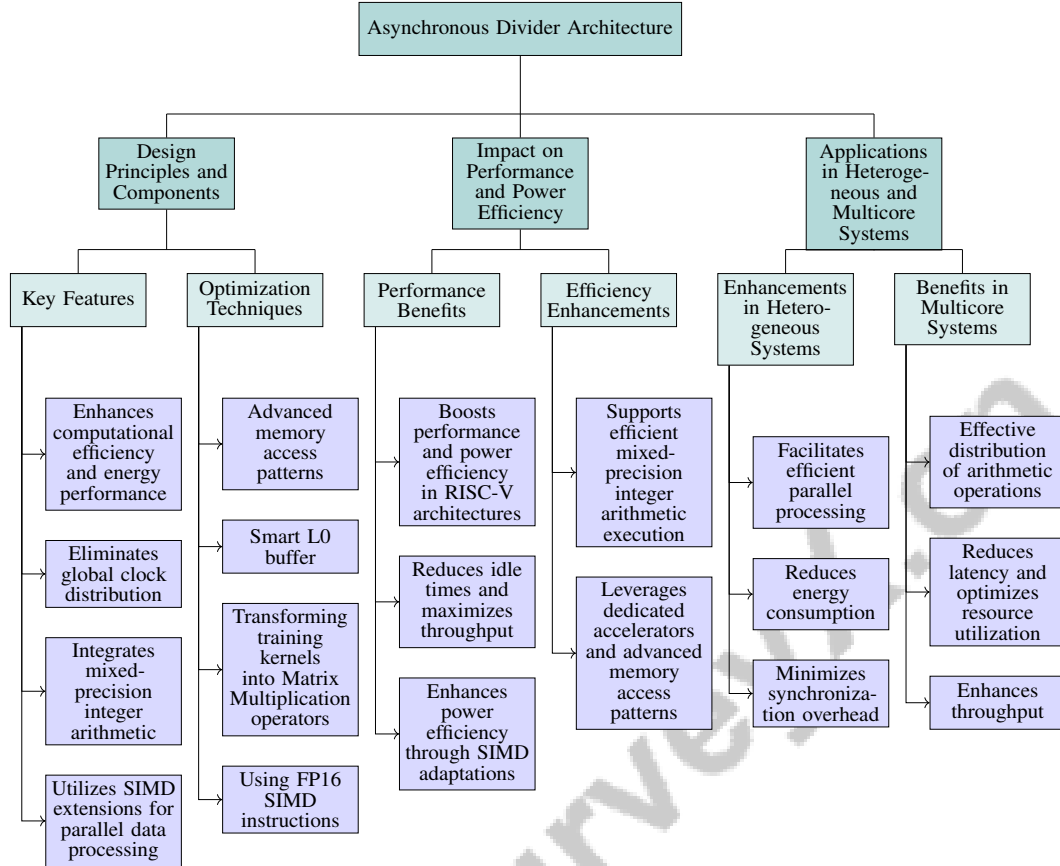


Figure 4: This figure illustrates the hierarchical structure of the Asynchronous Divider Architecture, detailing its design principles, impact on performance and power efficiency, and applications in heterogeneous and multicore systems. The architecture's key features and optimization techniques are highlighted, emphasizing its role in enhancing computational efficiency and energy performance. The figure also outlines the performance benefits and efficiency enhancements offered by asynchronous dividers, showcasing their adaptability and contributions to modern computing environments.

ALU operations, and data management, vital for improving processor performance in asynchronous divider architecture [15]. Additionally, virtual prototypes using SystemC and TLM-2.0 support early software development and system-level validation, ensuring robustness and adaptability of asynchronous divider architectures [7].

## 4.2 Impact on Performance and Power Efficiency

Integrating asynchronous dividers into computing systems significantly boosts performance and power efficiency, particularly within RISC-V processor architectures. By removing global clock distribution, asynchronous dividers allow computations and data transfers to occur independently, reducing idle times and maximizing throughput, leading to enhanced computational efficiency and energy savings [16]. The performance benefits of SIMD adaptations, especially in modular arithmetic, enhance power efficiency through parallel processing of multiple data elements [33]. For instance, the Nanhu-vdot processor achieves about 30

Moreover, asynchronous divider architecture supports efficient mixed-precision integer arithmetic execution, crucial for optimizing DNN processing. By leveraging dedicated accelerators and advanced memory access patterns, these dividers enhance RISC-V cores' performance in heterogeneous computing clusters, contributing to overall system efficiency [16].

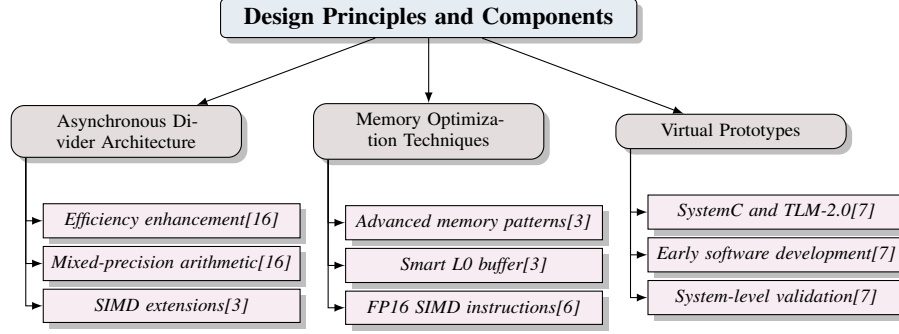


Figure 5: This figure illustrates the hierarchical structure of design principles and components in asynchronous divider architecture, memory optimization techniques, and virtual prototypes. It highlights key elements such as efficiency enhancement, mixed-precision arithmetic, SIMD extensions, and the role of SystemC and TLM-2.0 in early software development and system-level validation.

### 4.3 Adaptability to SIMD Architectures

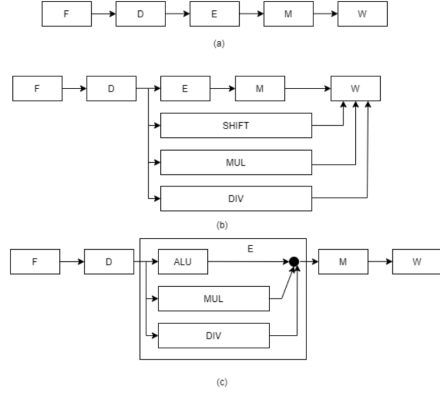
The adaptability of asynchronous dividers to SIMD architectures is pivotal for enhancing computational efficiency and energy performance in modern processors. Their capability to manage parallel processing tasks without synchronous clocking allows for simultaneous execution of multiple data operations, beneficial for applications demanding high throughput and energy efficiency [3]. Incorporating asynchronous dividers into SIMD architectures enables efficient execution of vectorized operations, essential for accelerating tasks like matrix multiplications and CNN operations. Transforming training kernels into Matrix Multiplication operators and applying FP16 SIMD instructions exemplify performance optimization potential [6]. By harnessing SIMD's parallel processing capabilities, asynchronous dividers can significantly lower computational latency and power consumption, making them ideal for resource-constrained environments such as IoT and edge computing [33].

Furthermore, the adaptability of asynchronous dividers facilitates developing specialized hardware accelerators that optimize arithmetic operations for improved performance, essential for supporting advanced applications, including AI and machine learning. Integrating asynchronous dividers into SIMD architectures aligns with broader digital design optimization goals, focusing on balancing computational performance with energy efficiency [16].

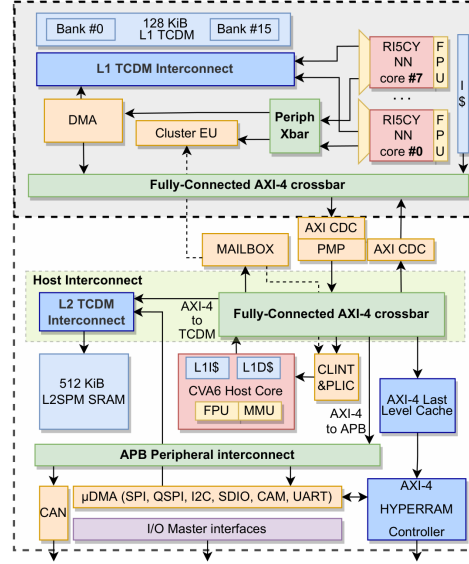
As illustrated in Figure 6, adaptability to SIMD architectures is crucial for optimizing computational efficiency and performance within computer architecture. The asynchronous divider architecture exemplifies this adaptability, as shown in the first subfigure, which presents three distinct block diagrams of a computer system, each reflecting varying complexity and integration levels. The initial diagram (a) outlines a basic flowchart with a single input and output, while the second diagram (b) features a more intricate system with multiple inputs and outputs, indicative of sophisticated data processing pathways. The second subfigure provides a detailed depiction of a computer system's interconnect and memory hierarchy, highlighting components such as L1 and L2 caches, HyperRAM controllers, and various interconnects like APB and AXI-4. These diagrams collectively underscore the flexibility and scalability of SIMD architectures in accommodating diverse computational demands, thereby enhancing the overall performance and efficiency of modern computer systems [20, 34].

### 4.4 Applications in Heterogeneous and Multicore Systems

The application of asynchronous dividers in heterogeneous and multicore systems significantly enhances computational capabilities by facilitating efficient parallel processing and reducing energy consumption. In heterogeneous systems, where diverse processors collaborate to execute complex tasks, asynchronous dividers enable non-blocking operations and minimize synchronization overhead, thereby boosting overall system performance. This capability is especially beneficial in environments with varied workloads, such as edge computing and IoT applications [16]. In multicore systems, asynchronous dividers allow for the effective distribution of arithmetic operations across multiple cores, enhancing throughput and reducing latency. By removing the need for global clock synchronization, each core can operate independently, optimizing resource utilization and power efficiency. This



(a) The image shows three different block diagrams of a computer system.[20]



(b) The image depicts a block diagram of a computer system's interconnect and memory hierarchy.[34]

Figure 6: Examples of Adaptability to SIMD Architectures

independence is critical for maintaining high performance in systems with fluctuating computational demands, particularly in AI and machine learning applications [3].

The adaptability of asynchronous dividers to heterogeneous and multicore architectures also fosters the development of specialized accelerators that enhance the execution of parallel tasks. These accelerators can be tailored to optimize specific operations, such as matrix multiplications and vector dot products, vital for high-performance computing in data-intensive applications. Integrating asynchronous dividers into these architectures aligns with broader digital design optimization goals, aiming to balance computational performance with energy efficiency [33]. Incorporating asynchronous dividers into heterogeneous and multicore systems strategically enhances computational capabilities by addressing challenges related to power efficiency, reliability, and performance scalability in modern architectures. This is particularly relevant in open-source RISC-V processor implementations and the increasing demand for specialized hardware in compute-intensive applications like artificial intelligence and edge computing [19, 18, 1]. By promoting efficient parallel processing and reducing energy consumption, asynchronous dividers contribute to the development of high-performance, low-power computing solutions that meet the demands of contemporary digital applications.

## 5 RISC-V Processor Framework

### 5.1 Open-Source Nature and Modular Design

The RISC-V processors' open-source and modular design fosters innovation and customization across diverse applications, allowing developers to tailor designs to specific needs, thereby enhancing computational efficiency and integrating advanced functionalities [2]. This flexibility is particularly advantageous in educational and resource-constrained settings where simplicity and minimal resource utilization are crucial [11]. The open-source ethos of RISC-V facilitates the development of specialized algorithms, as demonstrated by the FANN-on-MCU framework, which optimizes neural network deployment on RISC-V processors [2]. NoX, a compact open-source RISC-V core supporting FreeRTOS, exemplifies RISC-V's design principles by enhancing system integration and performance customization in Multi-Processor Systems-on-Chip (MPSoCs) [13].

---

Benchmarking initiatives, such as those by Dörflinger et al., emphasize the importance of standardized benchmarking for evaluating application-class RISC-V processors across multiple criteria, ensuring accurate performance assessment [21]. The open-source and modular design also supports the development of safety-critical applications, as seen in SentryCore, which enhances innovation and customization for robust security measures [14]. Optimization techniques for near-threshold voltage operations further illustrate RISC-V's potential for high energy efficiency [6]. Recent advancements in RISC-V architectures and specialized hardware for machine learning enable the creation of tailored architectures that address the diverse needs of modern technology, from educational settings to complex safety-critical systems in robotics, aerospace, and automotive sectors [17, 35, 1, 19].

## 5.2 Micro-Decoding Unit Integration

Integrating micro-decoding units within RISC-V architecture is crucial for optimizing processor performance by translating macro-instructions into micro-instructions, enhancing execution efficiency in complex computing environments [17]. This process allows granular control over processor operations, improving parallelism and resource utilization. The micro-decoding unit's contribution to RISC-V's modularity enables developers to implement customized instruction sets for specific applications, particularly in digital signal processing and machine learning. The dynamic micro-decoding unit significantly improves computational throughput and reduces latency, as exemplified by the Xiangshan Nanhu core, which achieves over four times the speed of traditional scalar methods for large language model inference tasks [17, 9].

Furthermore, micro-decoding units optimize pipeline architecture in RISC-V processors, reducing bottlenecks and enhancing efficiency in edge AI applications and large language model inference computations [17, 9]. This optimization is vital for high-performance applications requiring rapid processing of large datasets, such as real-time analytics or edge computing. The micro-decoding unit thus plays a critical role in RISC-V architecture, enhancing modularity and adaptability to meet varied computational requirements while improving security, energy efficiency, and execution speed [17, 19].

## 5.3 Instruction Set Extensions

Instruction set extensions are essential for enhancing RISC-V processors' capabilities, allowing for customization and performance optimization to meet specific application demands. These extensions facilitate incorporating specialized instructions that improve computational efficiency, broadening the scope of applications supported by RISC-V processors [2]. RISC-V's modular ISA enables seamless integration of custom extensions, vital for optimizing performance in machine learning, digital signal processing, and cryptography. Vector processing extensions, for instance, enhance data-parallel task handling, significantly boosting performance in applications requiring high throughput and low latency [3].

Moreover, instruction set extensions bolster RISC-V processors' security features by incorporating specialized instructions for secure data handling and encryption, ensuring data integrity and confidentiality in security-critical applications [14]. The flexibility to develop custom extensions also advances educational and research initiatives, allowing experimentation with novel computational models and architectural innovations [13].

## 5.4 Formal Verification and Validation

Formal verification and validation are critical for ensuring RISC-V processors' reliability and security, particularly given their open-source and modular nature. These processes involve rigorous analysis of processor designs to verify compliance with specified requirements and performance under all conditions, preventing errors and vulnerabilities that could compromise system integrity [14]. In RISC-V processors, formal verification is crucial for validating custom instruction set extensions and micro-architectural modifications, ensuring enhancements do not introduce unforeseen errors or vulnerabilities, especially in safety-critical applications [13].

Formal verification also supports implementing advanced security features by providing a framework for proving the correctness of cryptographic algorithms and secure communication protocols, essential for maintaining data integrity and confidentiality. Integrating formal verification techniques with

---

design automation tools enhances the efficiency of verification processes, allowing rapid validation of complex processor designs and reducing time-to-market [3]. These processes not only enhance processor design quality and robustness but also contribute to innovation and sustainability in digital design [29, 17, 19].

## **6 Low-Power Computing in RISC-V**

Enhancing energy efficiency and optimizing performance are central to advancing low-power computing. A key strategy involves integrating heterogeneous computing and specialized accelerators, which employ diverse processing units to enhance computational efficiency while minimizing power usage. This approach supports complex task execution and strategic workload distribution across different architectures. The following subsection explores the critical role of heterogeneous computing and specialized accelerators in low-power computing advancement.

### **6.1 Heterogeneous Computing and Specialized Accelerators**

Heterogeneous computing and specialized accelerators are crucial for optimizing performance and energy efficiency in low-power computing. Integrating RISC-V processors with ARM multi-core systems, as seen in mixed-criticality multi-OS architectures, allows workload balancing across processors, reducing power consumption and enhancing computational efficiency [36]. The HULK-V platform exemplifies the benefits of heterogeneous computing in low-power settings, offering high energy efficiency and the capability to run a full Linux OS on cost-effective SoC architectures, ideal for embedded and IoT applications [34].

Specialized accelerators further enhance these systems by providing tailored solutions for specific tasks. The integration of M-extension instructions into soft processors enhances integer arithmetic operations in RISC-V processors, reflecting both practical and theoretical performance improvements [20]. These accelerators are optimized for applications like DNN processing, where mixed-precision formats and advanced memory access patterns are crucial for low-power operation [16]. Insights into developing specialized accelerators can be drawn from quantum circuits for multilevel qudits, enhancing computational efficiency in tasks requiring high precision and low latency [28].

The modularity and standardized protocols in heterogeneous multicore SoCs facilitate the integration of various RISC-V cores and accelerators, promoting the development of energy-efficient computing systems [18]. This flexibility is vital for adapting to emerging computational models and applications, ensuring heterogeneous computing systems meet modern technology demands.

### **6.2 Educational and Real-Time Applications**

Implementing low-power computing strategies in educational and real-time systems fosters innovation and practical learning while ensuring efficient resource use. Lightweight, open-source RISC-V processors are ideal for educational purposes, providing an intuitive platform for demonstrating efficient hardware utilization [11]. This approach simplifies learning for students and educators and offers a practical framework for exploring advanced computing concepts in resource-constrained environments.

In educational settings, RISC-V's open-source nature allows extensive customization and experimentation, enabling hands-on experience with processor design and optimization. This flexibility is crucial for developing a deep understanding of low-power computing strategies and their real-world applications. Educators can customize learning experiences to cover topics like instruction set architecture, SoC design, and specialized hardware integration for machine learning and AI applications [17, 18, 19, 20].

Real-time systems, requiring precise timing and reliable performance, benefit from low-power strategies. Efficient task execution while minimizing power consumption is critical for applications like autonomous vehicles and industrial automation. RISC-V processors provide a powerful platform for executing energy-efficient algorithms, enhancing computational processes to meet demanding real-time standards. The evolution of RISC-V architecture now prioritizes security, energy efficiency, and execution speed, akin to CISC architectures. Modular open-source platforms like ESP facilitate the development of heterogeneous multicore systems, optimizing system performance [17, 18].

---

The intersection of educational and real-time applications highlights the significance of low-power strategies in advancing academic and practical technological advancements. RISC-V processors serve as crucial educational tools, offering a versatile platform for exploring complex computing concepts and equipping future engineers and developers with the skills necessary to tackle modern computing challenges, including demands for enhanced security, energy efficiency, and support for heterogeneous, multicore SoC architectures in fields like AI and machine learning [17, 18, 19].

### 6.3 Dynamic Instruction Execution and Micro-Decoding

Dynamic instruction execution enhances RISC-V processors' power efficiency by optimizing instruction flow and reducing binary size, supporting post-fabrication updates essential for maintaining security and efficiency [17]. This adaptability ensures RISC-V processors remain competitive over time. Integrating micro-decoding units within the RISC-V architecture further enhances power efficiency by translating complex macro-instructions into micro-instruction sequences, allowing granular control over execution and improving parallelism and resource utilization [17].

Principles of quantum computing, such as superposition and entanglement, exemplified by the Quantum Multiply-Add Circuit (QMAC), provide insights into optimizing instruction execution. These principles enable simultaneous operations, reducing overall circuit depth and enhancing computational efficiency [22]. Adopting similar strategies can significantly improve RISC-V processors' power efficiency and performance, particularly for high-speed arithmetic operations.

### 6.4 Performance Improvements through Extensions and SIMD Adaptations

Performance improvements in low-power computing are driven by integrating instruction set extensions and SIMD adaptations within RISC-V processors. These enhancements facilitate vectorized operations, crucial for optimizing computational efficiency in applications like signal processing, machine learning, and linear algebra. Vectorized implementations of modular arithmetic outperform classical algorithms, contributing to substantial performance gains in low-power environments [33].

Specialized testing approaches optimize instruction execution, exemplifying potential performance improvements. A proposed method processes over 200 million instructions per hour, highlighting efficiency gains through targeted optimizations and high-performance testing frameworks [29]. Such methodologies ensure RISC-V processors meet modern computing applications' stringent demands while maintaining low power consumption.

Integrating SIMD adaptations in RISC-V-based systems facilitates parallel processing, essential for handling data-intensive tasks. An experimental setup with a 4x4 Processing Element (PE) array integrated into a RISC-V-based SoC with independent memory nodes underscores this approach's effectiveness. Evaluations using benchmarks from various domains demonstrate significant performance improvements achievable through such architectural innovations [8].

## 7 Integer Arithmetic and Digital Design Optimization

Integer arithmetic plays a pivotal role in optimizing digital design, particularly in enhancing energy efficiency and computational performance. This section explores various methodologies that underscore the importance of integer arithmetic in modern computing systems, with a focus on heterogeneous System-on-Chip (SoC) architectures and low-bitwidth SIMD (Single Instruction, Multiple Data) operations as key optimization strategies.

### 7.1 Heterogeneous SoCs and Low-Bitwidth SIMD Operations

In heterogeneous SoC architectures, low-bitwidth SIMD operations are crucial for improving energy efficiency and processing speed. These architectures integrate diverse processing units like CPUs, GPUs, and specialized accelerators to execute computational tasks efficiently. Leveraging low-bitwidth SIMD operations significantly enhances parallel processing capabilities, essential for data-intensive tasks in power-constrained environments [33]. These operations are particularly beneficial in machine learning and signal processing, where performance and energy savings can be prioritized over precision. By performing multiple operations concurrently on smaller data widths, computational load and power consumption are reduced, as demonstrated by modular arithmetic operations that



---

outperform traditional scalar methods [3, 33]. Integrating low-bitwidth SIMD operations into heterogeneous SoCs also facilitates the development of specialized hardware accelerators for parallel operations, such as matrix multiplications and vector dot products, aligning with digital design optimization goals [16].

## 7.2 Reversible Computing and Approximate Computing

Reversible and approximate computing paradigms enhance energy efficiency and computational performance. Reversible computing minimizes energy dissipation by allowing computational processes to be reversed without information loss, applicable in quantum computing and low-power embedded systems [26]. This approach adheres to thermodynamic reversibility principles, enabling computations with minimal energy loss through reversible logic gates. Techniques like reduced precision floating-point operations and mixed-precision computing are increasingly adopted in edge computing and IoT applications, optimizing computational efficiency while minimizing environmental impact [6, 1]. Approximate computing leverages the tolerance for imprecision in many applications, allowing deliberate approximations that significantly reduce energy usage and processing time. This paradigm is effective in multimedia processing, machine learning, and sensor data analysis, where perfect accuracy is often unnecessary. By relaxing precision, approximate computing alleviates computational burdens, making it suitable for resource-constrained environments like edge computing and IoT applications. Advancements in computer arithmetic, including fixed-point and integer arithmetic optimizations, further enhance its feasibility [6, 1, 4, 25, 27]. Integrating reversible and approximate computing techniques into digital design optimization yields substantial benefits in energy efficiency and performance enhancement, particularly in resource-constrained environments like edge computing and IoT, where innovative arithmetic techniques such as bfloat16 and posit arithmetic are explored [1, 26].

## 7.3 Integer Arithmetic Optimization and Energy Efficiency

Optimizing integer arithmetic is essential for enhancing energy efficiency in digital systems, especially in resource-constrained environments. The Quantum Multiply-Add Circuit (QMAC) reduces depth complexity, facilitating faster multiply-add operations and improving overall efficiency while lowering energy consumption [22]. The application of datatype defining rewrite systems (DDRSes) provides a robust framework for specifying arithmetic operations, ensuring efficient and reliable integer computations [23]. Techniques like LocalSMT, which focus on isolating Boolean or integer variables, improve the solving of satisfiability modulo theories (SMT) problems, optimizing integer arithmetic operations and contributing to energy efficiency [24].

## 7.4 Energy Efficiency in Edge and IoT Applications

Energy-efficient design is crucial in edge and IoT applications, which demand effective operation in resource-constrained environments. These applications require architectures that prioritize low power consumption while maintaining high performance, particularly for battery-powered devices. The integration of RISC-V processors offers a promising solution due to their open-source nature and modular design, facilitating the development of customized, energy-efficient architectures [2]. Processing data locally in edge computing minimizes communication with centralized servers, conserving energy and reducing latency. This capability is critical for IoT applications, supporting real-time data processing in smart cities, healthcare monitoring, and industrial automation. Low-power RISC-V cores enable efficient execution of complex algorithms, enhancing applications like machine learning and data analytics [3]. The FANN-on-MCU toolkit optimizes neural network deployment on RISC-V-based microcontrollers, demonstrating significant energy savings by tailoring architectures to low-power devices [2]. Specialized accelerators in heterogeneous computing systems further enhance computational efficiency in edge devices, enabling complex tasks with minimal energy expenditure [16]. Focusing on energy-efficient design maximizes device lifespan and sustains IoT ecosystems. By leveraging RISC-V architectures, developers can create tailored solutions that address the diverse demands of modern digital applications, supporting the integration of heterogeneous multicore systems essential for contemporary computing while aligning with sustainability and technological innovation objectives. Advancements like dynamic micro-decoding units and extensible virtual prototypes enhance RISC-V processors' capabilities, enabling efficient execution of multithreaded

---

applications and facilitating design space exploration for improved performance and power validation [17, 7, 18].

### 7.5 Quantum and Classical Methods in Integer Arithmetic

The integration of quantum and classical methods in integer arithmetic signifies a groundbreaking advancement in computational optimization. Hybrid quantum-classical circuits, such as the proposed quantum multiply-add circuit (QMAC) utilizing the Quantum Fourier Transform (QFT), demonstrate significant reductions in computational steps, requiring only  $O(n + k)$  for exact results and  $O(n + \log k)$  for approximations, compared to the classical approach demanding  $O(n \log k)$  steps. This advancement enhances performance and energy efficiency, addressing limitations in classical computational methods [22, 27, 28]. Quantum computing introduces techniques like QFT and QMAC, achieving considerable reductions in computational complexity and depth, facilitating efficient integer arithmetic through superposition and entanglement principles. These quantum circuits enable simultaneous operations, significantly reducing processing time relative to classical methods [22]. Classical methods also optimize integer arithmetic through advancements in algorithmic design and hardware implementation, with DDRSes providing a framework for efficient arithmetic operations, ensuring ground-completeness and resource efficiency [23]. Specialized accelerators, such as QMAC, optimize arithmetic operations for specific applications, leveraging Quantum Fourier Transform techniques to enhance integer arithmetic speed, achieving  $O(n + k)$  for exact results and  $O(n + \log k)$  for approximations, compared to  $O(n \log k)$  in traditional methods. This advancement paves the way for embedding quantum devices within conventional architectures, improving performance in fields like machine learning and edge computing [33, 22, 1, 28]. These accelerators can leverage quantum principles while remaining compatible with classical architectures, maximizing performance gains in energy-sensitive environments and aligning with digital design optimization objectives.

### 7.6 Adaptive Techniques in Integer Arithmetic

Adaptive techniques in integer arithmetic enhance computational flexibility and efficiency, particularly in performance-demanding environments. Refining datatype defining rewrite systems (DDRSes) provides a structured framework for specifying arithmetic operations, ensuring ground-completeness and efficiency in integer operations [23]. Future work could explore additional notational systems, such as hexadecimal, for more compact representations and faster computations. Adaptive algorithms dynamically adjust computational parameters based on task requirements, crucial for applications with varying data types and precision levels, such as machine learning and digital signal processing. Implementing adaptive algorithms optimizes real-time operations—addition, multiplication, and division—by leveraging mixed-precision and integer arithmetic, ensuring efficient resource utilization in edge computing and embedded systems [27, 1, 4]. Integrating adaptive techniques with hybrid quantum-classical circuits significantly enhances integer arithmetic performance. The QMAC, utilizing QFT, achieves faster computation times than traditional methods, requiring  $O(n + k)$  for exact results and  $O(n + \log k)$  for approximations, compared to  $O(n \log k)$  for classical circuits. This advancement is particularly relevant in edge computing and low-power architectures, where efficient arithmetic operations are critical for applications ranging from machine learning to IoT devices [22, 1, 28]. Quantum principles enable simultaneous execution of multiple arithmetic operations, reducing overall complexity and enhancing throughput. Combining quantum techniques with classical adaptive methods achieves a balance between efficiency and flexibility, catering to diverse modern digital application needs.

### 7.7 Security and Efficiency in Integer Arithmetic

Balancing security and efficiency in integer arithmetic operations is critical in modern computing system design. Integer arithmetic supports various applications, including cryptographic algorithms, error-correcting codes, and digital signal processing, where both security and performance are paramount. This arithmetic underpins advanced algorithms, such as modular integer operations and fast modular Fourier transforms, critical for efficient computations in contemporary software environments. Optimized algorithms for arbitrary-precision arithmetic enhance performance regardless of computer word size limitations [33, 27]. The challenge lies in optimizing arithmetic operations to ensure computational efficiency while maintaining robust security features. In cryptography, integer arithmetic is essential for implementing secure algorithms, underpinning operations like modular

---

arithmetic and efficient computation techniques used in error-correcting codes and probabilistic modeling, ensuring integrity and confidentiality [33, 27, 25]. Secure arithmetic often requires additional computational resources, impacting performance. Developing techniques that minimize security-related overhead while preserving data integrity is crucial, especially for resource-constrained environments like IoT devices and embedded systems. Advanced computational models involving quantum circuits offer promising solutions for enhancing both security and efficiency in integer arithmetic. Quantum principles enable complex arithmetic operations with reduced depth complexity, improving throughput and energy efficiency [22]. Leveraging quantum techniques allows for designing secure arithmetic operations without compromising performance, aligning with sustainable and efficient digital design objectives. Incorporating adaptive techniques in integer arithmetic can significantly bolster security by dynamically modifying computational parameters to meet specific task demands. This adaptability is essential in applications where performance and security must be balanced, such as edge computing and probabilistic programming. Utilizing advancements in computer arithmetic, including high-precision and mixed-precision formats, systems can optimize arithmetic operations to enhance both efficiency and security in real-time processing environments [33, 1, 4, 25, 27]. This adaptability enables real-time optimization of arithmetic operations, ensuring efficient application of security measures without unnecessary overhead. Developing specialized algorithms that balance security and efficiency is essential for supporting advanced applications like machine learning and data analytics, where performance and data protection are critical.

## 8 Conclusion

The incorporation of no-restoring algorithms and asynchronous dividers within RISC-V processors represents a significant leap forward in the realm of low-power computing. These innovations enhance computational efficiency while minimizing energy consumption, which is crucial in energy-sensitive environments. The no-restoring algorithm streamlines division operations by reducing cycle counts and power usage, making it especially beneficial in constrained settings. Asynchronous dividers complement this by facilitating non-blocking operations, thereby improving processing speed and efficiency, particularly in heterogeneous and multicore systems.

RISC-V's open-source and modular architecture enables extensive customization, fostering the development of specialized instructions and extensions that boost performance across a variety of applications, including those in AI and machine learning. SIMD adaptations further enhance computational efficiency through parallel processing, which is essential for optimizing modular arithmetic across diverse architectures.

Future research should focus on exploring additional RISC-V extensions and their integration into soft processor architectures, as well as optimizing compiler strategies to improve performance. The flexibility of the micro-decoding unit and the development of specialized instructions for security purposes present promising areas for further investigation. Moreover, refining quantization processes and examining ISA extensions to support lower precision data types could significantly advance methodologies in low-power computing.

The findings underscore RISC-V's potential in next-generation applications, offering substantial opportunities for hardware optimizations and affirming its capability to support complex software stacks. Additionally, the development of sound reversible programming languages that effectively manage side effects presents an exciting avenue for future research. Therefore, the strategic integration of no-restoring algorithms and asynchronous dividers in RISC-V processors is a promising approach to achieving high-performance, energy-efficient computing, with numerous opportunities for future exploration and advancement.

---

## References

- [1] Andre Guntoro, Cecilia De La Parra, Farhad Merchant, Florent De Dinechin, John L Gustafson, Martin Langhammer, Rainer Leupers, and Sangeeth Nambiar. Next generation arithmetic for edge computing. In *2020 Design, Automation & Test in Europe Conference & Exhibition (DATE)*, pages 1357–1365. IEEE, 2020.
- [2] Xiaying Wang, Michele Magno, Lukas Cavigelli, and Luca Benini. Fann-on-mcu: An open-source toolkit for energy-efficient neural network inference at the edge of the internet of things, 2022.
- [3] Michael Gautschi, Pasquale Davide Schiavone, Andreas Traber, Igor Loi, Antonio Pullini, Davide Rossi, Eric Flamand, Frank K. Gurkaynak, and Luca Benini. A near-threshold risc-v core with dsp extensions for scalable iot endpoint devices, 2016.
- [4] Takeshi Iwashita, Kengo Suzuki, and Takeshi Fukaya. An integer arithmetic-based sparse linear solver using a gmres method and iterative refinement, 2021.
- [5] Matěj Hejda, Federico Marchesin, George Papadimitriou, Dimitris Gizopoulos, Benoît Charbonnier, Régis Orobitchouk, Peter Bienstman, Thomas Van Vaerenbergh, and Fabio Pavanello. Invited: Neuromorphic architectures based on augmented silicon photonics platforms, 2024.
- [6] Reduced precision floating-point.
- [7] Vladimir Herdt, Daniel Große, Hoang M Le, and Rolf Drechsler. Extensible and configurable risc-v based virtual prototype. In *2018 Forum on Specification & Design Languages (FDL)*, pages 5–16. IEEE, 2018.
- [8] Daniel Vazquez, Jose Miranda, Alfonso Rodriguez, Andres Otero, Pascuale Davide Schiavone, and David Atienza. Strela: Streaming elastic cgra accelerator for embedded systems, 2024.
- [9] Xu-Hao Chen, Si-Peng Hu, Hong-Chao Liu, Bo-Ran Liu, Dan Tang, and Di Zhao. Research on llm acceleration using the high-performance risc-v processor "xiangshan" (nanhu version) based on the open-source matrix instruction set extension (vector dot product), 2024.
- [10] Michael Gautschi, Pasquale Davide Schiavone, Andreas Traber, Igor Loi, Antonio Pullini, Davide Rossi, Eric Flamand, Frank K. Gurkaynak, and Luca Benini. Near-threshold risc-v core with dsp extensions for scalable iot endpoint devices. *IEEE transactions on very large scale integration (VLSI) systems*, 25(10):2700–2713, 2017.
- [11] Ludovico Poli, Sangeet Saha, Xiaojun Zhai, and Klaus D McDonald-Maier. Design and implementation of a risc v processor on fpga. In *2021 17th international conference on mobility, sensing and networking (MSN)*, pages 161–166. IEEE, 2021.
- [12] Avani Dave, Nilanjan Banerjee, and Chintan Patel. Sracare: Secure remote attestation with code authentication and resilience engine, 2021.
- [13] Anderson I. Silva, Altamiro Susin, Fernanda L. Kastensmidt, Antonio Carlos S. Beck, and Jose Rodrigo Azambuja. Nox: a compact open-source risc-v processor for multi-processor systems-on-chip, 2024.
- [14] Michael Rogenmoser, Alessandro Ottaviano, Thomas Benz, Robert Balas, Matteo Perotti, Angelo Garofalo, and Luca Benini. Sentrycore: A risc-v co-processor system for safe, real-time control applications, 2024.
- [15] Hiromu Miyazaki, Takuto Kanamori, Md Ashraful Islam, and Kenji Kise. Rvcoprep: An optimized risc-v soft processor of five-stage pipelining. *IEICE TRANSACTIONS on Information and Systems*, 103(12):2494–2503, 2020.
- [16] Angelo Garofalo, Yvan Tortorella, Matteo Perotti, Luca Valente, Alessandro Nadalini, Luca Benini, Davide Rossi, and Francesco Conti. Darkside: A heterogeneous risc-v compute cluster for extreme-edge on-chip dnn inference and training, 2023.
- [17] Juliette Pottier, Thomas Nieddu, Bertrand Le Gal, Sébastien Pillement, and Maria Méndez Real. Risc-v processor enhanced with a dynamic micro-decoder unit, 2024.

- 
- [18] Joseph Zuckerman, Paolo Mantovani, Davide Giri, and Luca P. Carloni. Enabling heterogeneous, multicore soc research with risc-v and esp, 2022.
- [19] Stavros Kalapothas, Manolis Galetakis, Georgios Flamis, Fotis Plessas, and Paris Kitsos. A survey on risc-v-based machine learning ecosystem. *Information*, 14(2):64, 2023.
- [20] Md Ashraful Islam, Hiromu Miyazaki, and Kenji Kise. Rvcorep-32im: An effective architecture to implement mul/div instructions for five stage risc-v soft processors, 2020.
- [21] Alexander Dörflinger, Mark Albers, Benedikt Kleinbeck, Yejun Guan, Harald Michalik, Raphael Klink, Christopher Blochwitz, Anouar Nechi, and Mladen Berekovic. A comparative survey of open-source application-class risc-v processor implementations. In *Proceedings of the 18th ACM international conference on computing frontiers*, pages 12–20, 2021.
- [22] Christopher M. Maynard and Einar Pius. Integer arithmetic with hybrid quantum-classical circuits, 2013.
- [23] Jan A. Bergstra and Alban Ponse. Three datatype defining rewrite systems for datatypes of integers each extending a datatype of naturals, 2016.
- [24] Shaowei Cai, Bohan Li, and Xindi Zhang. Local search for satisfiability modulo integer arithmetic theories, 2023.
- [25] William X. Cao, Poorva Garg, Ryan Tjoa, Steven Holtzen, Todd Millstein, and Guy Van den Broeck. Scaling integer arithmetic in probabilistic programs, 2023.
- [26] Chris Heunen, Robin Kaarsgaard, and Martti Karvonen. Reversible effects as inverse arrows, 2018.
- [27] Richard P. Brent and Paul Zimmermann. Modern computer arithmetic (version 0.5.1), 2010.
- [28] Archimedes Pavlidis and Emmanuel Floratos. Arithmetic circuits for multilevel qudits based on quantum fourier transform, 2017.
- [29] Vladimir Herdt, Daniel Große, Eyck Jentzsch, and Rolf Drechsler. Efficient cross-level testing for processor verification: A risc-v case-study. In *2020 Forum for Specification and Design Languages (FDL)*, pages 1–7. IEEE, 2020.
- [30] Hengrui Zhao, Dong Liu, and Houqiang Li. Efficient integer-arithmetic-only convolutional neural networks, 2020.
- [31] Alex Parent, Martin Roetteler, and Michele Mosca. Improved reversible and quantum circuits for karatsuba-based integer multiplication. *arXiv preprint arXiv:1706.03419*, 2017.
- [32] Karthik Ganesan, Florian Lonsing, Srinivasa Shashank Nuthakki, Eshan Singh, Mohammad Rahmani Fadiheh, Wolfgang Kunz, Dominik Stoffel, Clark Barrett, and Subhasish Mitra. Effective pre-silicon verification of processor cores by breaking the bounds of symbolic quick error detection, 2021.
- [33] Joris van der Hoeven, Grégoire Lecerf, and Guillaume Quintin. Modular simd arithmetic in mathemagix, 2014.
- [34] Luca Valente, Yvan Tortorella, Mattia Sinigaglia, Giuseppe Tagliavini, Alessandro Capotondi, Luca Benini, and Davide Rossi. Hulk-v: a heterogeneous ultra-low-power linux capable risc-v soc, 2022.
- [35] Juan-David Guerrero-Balaguera, Josie E. Rodriguez Condia, and Matteo Sonza Reorda. A novel compaction approach for sbst test programs, 2021.
- [36] Luca Cuomo, Claudio Scordino, Alessandro Ottaviano, Nils Wistoff, Robert Balas, Luca Benini, Errico Guidieri, and Ida Maria Savino. Towards a risc-v open platform for next-generation automotive ecus, 2023.

---

**Disclaimer:**

SurveyX is an AI-powered system designed to automate the generation of surveys. While it aims to produce high-quality, coherent, and comprehensive surveys with accurate citations, the final output is derived from the AI's synthesis of pre-processed materials, which may contain limitations or inaccuracies. As such, the generated content should not be used for academic publication or formal submissions and must be independently reviewed and verified. The developers of SurveyX do not assume responsibility for any errors or consequences arising from the use of the generated surveys.

www.SurveyX.cn